

Relational Model

Key Constraints

- **Superkey** \rightarrow subset of attributes that *uniquely* identifies a tuple in a relation. Attributes A is not a superkey if there are two tuples t_1 and t_2 such that
 - $\forall a \in A : t_1.a = t_2.a \wedge \exists b \in R - A : t_1.b \neq t_2.b$
- **Key** \rightarrow superkey that is **minimal**
- **Primary key** \rightarrow selected candidate key for a relation
 - *must* be non-NULL and UNIQUE (**entity integrity constraint**)
 - Candidate keys may be NULL or contain NULL columns.

Foreign Key Constraints

- A **foreign key** is a subset of attributes of relation $R1$ that
 - 1. appear as a **primary key** in the referenced relation $R2$
 - 2. be a NULL value (OR contain at least one NULL value)

Relational Algebra

Relations are a **set**. Relations are closed under Relational Algebra.

Conjunction					Disjunction				
$c_1 \wedge c_2$		c_1			$c_1 \vee c_2$		c_1		
		False	NULL	True			False	NULL	True
c_2	False	False	False	False		False	False	NULL	True
	NULL	False	NULL	NULL	c_2	NULL	NULL	NULL	True
	True	False	NULL	True		True	True	True	True

Binary Operators

Set Operators

Two relations are **union-compatible** if

- R and S have the same number of attributes;
- corresponding attributes have the *same or compatible domains*; and no requirement for same attribute names

Cross product

$R \times S$ returns a relation with schema (A, B, X, Y) where $R \times S = \{(a, b, x, y) \mid (a, b) \in R, (x, y) \in S\}$

- $|R \times S| = |R| * |S|$ (note **empty tables**)
- Attributes A and B **must be disjoint**

Division

Consider two relations $R(A_1, \dots, A_m), S(B_1, \dots, B_n)$, where $attr(S) \subset attr(R)$. The division of R by S computes the largest set of tuples $Q \subset \pi_{[A_1, \dots, A_m]}(R)$ such that for every tuple $(a_1, \dots, a_m) \in Q$,

$$\pi_{[L]}((a_1, \dots, a_m) \times S) \in R$$

R/S is equivalent to $\pi_{[A]}(R) - \pi_{[A]}((\pi_A(R) \times S) - R)$

Join operators

Inner Joins

- **Natural Join** \bowtie Let A be the set of attributes that R and S have in common, $c = \forall a_i \in A : R_{a_i} = S_{a_i}$ and $\ell = Attr(R) \cup (Attr(S) - A)$, then

$$R \bowtie S = \pi_{\ell}(R \bowtie_c S)$$

Outer Joins

- **dangle** $(R \bowtie_{\theta} S) \rightarrow$ set of dangling tuples in R wrt to $R \bowtie_{\theta} S - dangle(R \bowtie_{\theta} S) \subseteq R$
- **null** $(R) \rightarrow n$ -component **tuple** of null values where n is the number of attributes of R

Types

- **Left outer join**

$$R_{\theta} S = R \bowtie_{\theta} S \cup dangle(R \bowtie_{\theta} S) \times \{null(S)\}$$

- **Right outer join**

$$R_{\theta} S = R \bowtie_{\theta} S \cup (\{null(R)\} \times dangle(S \bowtie_{\theta} R))$$

- **Full outer join**

$$R_{\theta} S = R \bowtie_{\theta} S \cup (dangle(R \bowtie_{\theta} S) \times \{null(S)\})$$

$$\cup (\{null(R)\} \times dangle(S \bowtie_{\theta} R))$$

- **Natural Outer join** is analogous to natural inner join where equality operation is performed over all attributes that R and S have in common.

Equivalence

- **Strongly** equivalent if for any input **both produces error or both produces the same result**.
- **Weakly** equivalent if for any input there is no error then both produces the same result.

- $\sigma_{[c]}(A \cup B) \neq \sigma_{[c]}(A) \cup \sigma_{[c]}(B)$ (invalid columns).
- $\pi_{[A]}(R - S) \neq \pi_A(R) - \pi_A(S)$ (e.g. 12, 13).
- $\pi_{[D,Y]}(R \times S) \equiv \pi_{[D,Y]}(S \times R) \equiv \pi_{[D]}(S) \times \pi_{[Y]}(R)$
- $R \bowtie_{[\theta_1]} (S \bowtie_{[\theta_2]} T) \neq (R \bowtie_{[\theta_1]} S) \bowtie_{[\theta_2]} T$ unless θ_1 uses T and θ_2 uses R.
- $(E_1 \bowtie_{c_1} E_2) \bowtie_{c_2 \wedge c_3} E_3 \neq E_1 \bowtie_{c_1 \wedge c_2} (E_2 \bowtie_{c_3} E_3)$
- $\sigma_{[\theta]}(R \times S) \neq \sigma_{[\theta]}(R) \times S$

SQL Overview

Integrity Constraints

Unique constraint

The following expression should hold at all times:

$$\forall r_1, r_2 : (r_1 \equiv r_2) \vee (\exists a : r_1.a <> r_2.a)$$

Since the *unique* check is done via $<>$, as long as the tuple in question contains a NULL value, then the condition always evaluates to NOT false.

Deferrable Constraints

Constraints are checked immediately at the end of a SQL statement. This is true even for **transactions**.

- Available for UNIQUE, PRIMARY KEY, FOREIGN KEY.

```
ALTER TABLE <TBL>
ADD CONSTRAINT <NAME>
FOREIGN KEY(Attr) REFERENCES TBL(Attr)
DEFERRABLE INITIALLY DEFERRED | IMMEDIATE;
```

Considerations

- No need to care about the order of statements inside a transaction
- Allows for cyclic foreign key constraints
- Data definition no longer unambiguous
- Certain checks need to be done at run-time (performance penalty).

Querying the database

By default, PostgreSQL allows duplicate tuples to exist, and row-ordering is not guaranteed (order independent).

Common SQL Constructs

Grouping


Given GROUP BY a_1, \dots two tuple t and t' belong to the same group if $r1.a1 \equiv r2.a1, \dots$ If column A_i of **table R** appears in the SELECT/HAVING clause, either

1. A_i appears in the GROUP BY clause
2. A_i appears as input of an aggregation function in the SELECT/HAVING clause
3. The primary key of R appears in the GROUP BY clause

Remark:

1. There should be unique values for each group. If not, then the result is undefined.
2. Group based on **equivalence** (incl. NULL)
3. If the SELECT clause has multiple attributes from different tables, the same condition applies to each individual attribute.

Conceptual Evaluation of Queries



Conditional expressions

case statements: Very similar to how a normal switch-case statement is written.

```
-- preferred
CASE [expression]
WHEN <condition> | <value> THEN <result>
WHEN <condition> | <value> THEN <result>
ELSE <default>
END
```

Structuring SQL Queries

Views

Only parts of the table is of interest. Equivalent to an external schema that is presented to the user.

Note:

- A VIEW is a permanently named query. The computation may be done each time the virtual table is accessed (i.e. not permanently stored).
- Can be used as a normal table.
- For INSERT, UPDATE, and/or DELETE statements.
 1. Only one entry in the FROM clause
 2. No WITH DISTINCT, GROUP BY, LIMIT or OFFSET
 3. No UNION, INTERSECT, EXCEPT or ALL
 4. No aggregates

Extended Concepts

Recursive Queries

```
WITH RECURSIVE Linker(to_stn, stops) AS (
  SELECT to_stn, 0
  FROM MRT WHERE fr_stn = "NS1"
  UNION [ ALL ]
  SELECT M.to_stn, L.stops + 1
  FROM Linker L, MRT M
  WHERE M.fr_stn = L.to_stn
  AND L.stops < 3 -- termination
)
SELECT to_stn, stops FROM Linker;
```

General Guidelines

Relation cardinality

Suppose entities A and B has a relationship R. The following table shows the cardinality of the relationship R.

Relationship	Min	Max
(total)-to-(many)	A	AB
(total)-to-(key)	A	B
\rightarrow error if $A > B$		
(key)-to-(many)	0	A
(total+key)-to-(many)	A	A
(total)-to-(total)	max(A, B)	AB
(key)-to-(key)	0	min(A, B)
(many)-to-(many)	0	AB
(total+key)-to-(total)	A	A
\rightarrow error if $A < B$		
(total+key)-to-(key)	A	A
\rightarrow error if $A > B$		
(total+key)-to-(total+key)	A=B	A=B
\rightarrow error if $A <> B$		

Procedures and Functions

Statment level interfaces allows us to execute SQL in a host language. **Call level interfaces** allows us to code in the host language only.

Functions

Functions are reusable snippets that **return some value**. General syntax:

```
CREATE OR REPLACE FUNCTION
<NAME>(<<IN|OUT|INOUT> ARGS <TYPE>)>
RETURNS <SQL_RETURN_TYPE | [SETOF] TBL_NAME |
RECORD | SETOF RECORD> AS $$
:
$$ LANGUAGE sql
```

Procedures


Procedures are reusable snippets that **do not return any value** i.e. update. They are transactional.

General syntax:

```
CREATE OR REPLACE PROCEDURE <NAME>(<ARGS <TYPE>)>
AS $$
:
$$ LANGUAGE sql;
```

Cursors

Cursor workflow:



For example:

```
CREATE OR REPLACE FUNCTION score_gap()
RETURNS TABLE ( name TEXT, mark INT, gap INT )
AS $$ DECLARE
  -- declares a cursor variable
  curs CURSOR FOR (SELECT * FROM Scores ORDER
    BY Mark DESC);
  r RECORD;
  prv_mark INT;
BEGIN
  prv_mark := -1;
  OPEN curs;
  LOOP
    FETCH curs INTO r; -- read the input into r
    EXIT WHEN NOT FOUND; -- terminate
    name := r.Name;
    mark := r.Mark;
    IF prv_mark >= 0 THEN
      gap := prv_mark - mark;
    ELSE gap := NULL;
    END IF;
    RETURN NEXT; -- inserts tuple to output
    prv_mark := r.Mark;
  END LOOP;
  CLOSE curs; -- releases resources
END;
$$ LANGUAGE plpgsql;

FETCH -- some variants
[PRIOR | FIRST | LAST | ABSOLUTE NUM]
FROM cur INTO r
```

SQL Injection Attack

Use prepared statements instead of directly executing SQL.

```
-- instead of using this
EXEC SQL EXECUTE IMMEDIATE :stmt;

-- use this
EXEC SQL PREPARE mystmt FROM :stmt;
EXEC SQL EXECUTE mystmt USING :in_name;
```

Triggers

```
-- Trigger
CREATE TRIGGER <NAME>
<AFTER | BEFORE | INSTEAD OF> <INSERT | DELETE
  | UPDATE> ON <TBL>
[ WHEN <condition>]
FOR EACH <ROW | STATEMENT> EXECUTE FUNCTION <
  trigger_func>();

-- Trigger function
CREATE OR REPLACE FUNCTION <NAME>()
RETURNS TRIGGER AS $$
BEGIN
  :
END;
$$ LANGUAGE plpgsql;
```

A trigger function has access to a number of variables.

- NEW — the new tuple

- OLD — the old tuple (update/delete)
- TG.OP — the operation that activates the trigger
- TG.TABLE_NAME — name of table that trigger the invocation
- CURRENT_DATE — the current date

Notes

- Returning NULL in a BEFORE trigger tells the database to ignore the rest of the operation.
- OLD is set to NULL in a BEFORE INSERT trigger.
- Whenever the function returns a non-null tuple, the function will use it as the tuple to be inserted (incl. OLD).
- For BEFORE trigger, returning a non-null value will cause the operation to proceed as normal. Otherwise, no tuples will be inserted.
- The return value for an AFTER trigger is ignored.
- For INSTEAD OF trigger, returning non-null signals to proceed as normal, whereas null indicates to ignore the rest of the operations on the **current** row.
- Instead Of triggers can only be used on views.

```
CREATE TRIGGER update_max_trigger
INSTEAD OF UPDATE ON Max_Score -- view
FOR EACH ROW EXECUTE FUNCTION update_max_func();

CREATE OR REPLACE FUNCTION update_max_func()
RETURNS TRIGGER AS $$ BEGIN
  -- actual table
  UPDATE Scores SET Mark = NEW.Mark WHERE
    Name = OLD.Name;
  RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

Trigger levels

Statement-level triggers ignore the values returned by the trigger function. To omit subsequent operations, raise an **exception**.

Trigger timing

INSTEAD OF is only allowed on row-level.

Trigger condition

Condition in WHEN() can be more complicated subject to the following restrictions:

- No SELECT in WHEN()
- No OLD for INSERT
- No NEW for DELETE
- No WHEN for INSTEAD OF

Deferred Trigger

Defer the checking of triggers until the end of the transaction.

```
CREATE CONSTRAINT TRIGGER bal_check_trigger
AFTER INSERT OR UPDATE OR DELETE ON Account
DEFERRABLE INITIALLY DEFERRED
-- must be for each row
FOR EACH ROW EXECUTE FUNCTION bal_check_func();

-- if set INITIALLY IMMEDIATE
BEGIN TRANSACTION;
SET CONSTRAINTS bal_check_trigger DEFERRED;
:
COMMIT;
```

Order of triggers

Within each category, triggers are activated in alphabetic order. If a BEFORE row-level trigger returns NULL, then subsequent triggers on the same row are omitted.

- BEFORE statement-level
- BEFORE row-level
- AFTER row-level
- AFTER statement-level

Functional Dependencies

Two attributes are functionally equivalent if their closures are equal. All candidate keys are functionally equivalent.

- $\alpha \rightarrow \beta$ is a **trivial** FD if $\beta \subset \alpha$.
- $\alpha \rightarrow \beta$ is a **non-trivial** FD if $\beta \not\subset \alpha$.
- $\alpha \rightarrow \beta$ is a **completely non-trivial** FD if $\beta \cap \alpha = \emptyset$. Both A and B cannot be empty sets.

Computing attribute closures

A **closure** is a set of all attributes that are functionally dependent on a given set of attributes.

- Check if any attribute doesn't appear in the RHS of any FD. These attributes must appear in the key
- Compute attribute closure starting with singular attributes.
- Note all candidate keys in the process
- If current set of attributes is a superset of some previously seen, candidate key, can skip

Anomalies

- Insertion Anomaly** — Unable to insert a record into a table because the table's structure does not allow us to enter certain information.
- Update Anomaly** — Update one instance of the data without updating all other instances.
- Deletion Anomaly** — Loss of some, or all, of the information related to the deleted record.

Armstrong Axioms

Sound: The rule only generates elements of Σ^+ when applied to Σ **Complete:** The rule(s) generate(s) all elements of Σ^+ when applied to Σ

- $XY \rightarrow Y$ (Reflexivity)
- $X \rightarrow Y \Rightarrow XZ \rightarrow YZ$ (Augmentation)
- if $X \rightarrow Y$ and $Y \rightarrow Z \Rightarrow X \rightarrow Z$ (Transitivity)
- if $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$ (Union)
- if $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$ (Decomposition)
- if $A \rightarrow B$ and $BC \rightarrow D$, then $AC \rightarrow D$ (Pseudo-transitivity)
- if $A \rightarrow B$ and $C \rightarrow D$, then $AC \rightarrow BD$ (Composition)
- if $AC \rightarrow BC$, then $A \not\rightarrow B$ (Reverse augmentation)
- if $X \rightarrow Y$, then $XZ \rightarrow Y$ (Weak augmentation)

Normalization

BCNF

A table R is in BCNF, if every **non-trivial** and **decomposed** FD has a superkey as its left hand side.

In general, a violation occurs, iff we have a non-trivial closure that contains **more** attributes than the left hand side of the FD, but does not contain all the attributes in the table.

BCNF Algorithm

- Find a subset X of the attributes in R s.t. its closure is a violation of BCNF.

- Decompose R into two tables R_1 and R_2 s.t. R_1 contains $\{X\}^+$ and R_2 contains $(R - \{X\}^+) \cup X$.
- If R_1 or R_2 is not in BCNF, decompose them further.

Lossless join

The original table can always be reconstructed from the decomposed tables, whenever the common attributes in R_1 and R_2 constitutes a superkey in R_1 **or** R_2 .

- A binary decomp is lossless-join \iff full outer natural join of its two fragments equal the initial table. Otherwise it is lossy
- A binary decomp of R into R_1 and R_2 is lossless-join if $R = R_1 \cup R_2$ and either $R_1 \cap R_2 \rightarrow R_1$ or $R_1 \cap R_2 \rightarrow R_2$
- A decomp is lossless-join if there exists a sequence of binary lossless-join decomp that generates that decomp

Dependency preserving

A decomp of R with Σ into R_1, R_2, \dots, R_n with respective projected FDs $\Sigma_1, \Sigma_2, \dots, \Sigma_n$ is dependency preserving $\iff \Sigma^+ = (\Sigma_1 \cup \dots \cup \Sigma_n)^+$

- Check that the minimal cover are equivalent

Notes

- The BCNF decomposition of a table may not be unique.
- If a table has only 2 attributes, then it must be in BCNF.
- Each decomposition step gets rid of at least one BCNF violation \implies termination is guaranteed.
- If there are no FDs in a table, the table is already in BCNF \implies key of the table consists of all the attributes in the table.
- No update, deletion and insertion anomalies.
- Cannot derive original FDs from R_1 and $R_2 \implies$ dependencies may not be preserved (not equivalent).

3NF

A table is in 3NF iff for every non-trivial and decomposed FD either (1) the left hand side is a superkey, or (2) right hand side is a proper subset of a candidate key. 3NF is not as strict as BCNF \implies higher levels of redundancy.

3NF Decomposition Algorithm

Given: A table R , and a set S of FDs

- Step 1: Derive a minimal basis of S
- Step 2: In the minimal basis, combine the FDs whose left hand sides are the same
- Step 3: Create a table for each FD remained
- Step 4: If none of the tables contains a key of the original table R , create a table that contains a key of R (To ensure lossless join decomposition)
- Step 5: Remove redundant tables

The **minimal basis** of S is a simplified version of S , such that (1) every FD in S can be derived from M , (2) every FDs in M is non-trivial and decomposed FD, (3) no FD in M is redundant and (4) none of the attributes on the LHS are redundant.

Algorithm for Minimal Basis

- Transform the FDs, so that each right hand side contains only one attribute.
- Remove redundant attributes on the left hand side of each FD.
- Remove redundant FDs.

May have update and delete anomalies in rare cases.